Anshuman Mishra

anshumanmishra.v.2.8@gmail.com +91-7398296725 linkedin.com/in/anshuman-mishra-726781364

SUMMARY

Generative AI Engineer specializing in agentic AI, RAG, and LLM-powered systems, backed by Artificial Intelligence research experience (ASR/TTS/OCR) at Bhashini, India's National Language Translation Mission. Built a LangGraph-based multi-agent RAG platform (A.R.T.E.M.I.S) and an LLM voice-agent system (AI_Concierge), combined with REST API and backend engineering experience shipping B2G exam-platform features at scale. Comfortable taking an AI use case from architecture through POC, evaluation, and production deployment.

EXPERIENCE

Assistant Manager - Full Stack Developer & Product

Innovatiview India Ltd (B2G), Noida **April 2026 - Present**

- Designed and shipped REST APIs for identity verification (Aadhaar checks, document uploads, liveness flows) with human-in-the-loop approval workflows for exception handling, integrated with enterprise government databases and consumed by Vue.js and Flutter clients.
- Built a JSON-configurable form engine, replacing manual backend changes with template-driven automation across 20+ client registration flows.
- Developed and tuned a constraint-based candidate-to-center allocation algorithm balancing accessibility and travel-distance rules.

AI/ML Intern

Bhashini, Ministry of Electronics and Information Technology **August 2025 - November 2025**

- Engineered a multi-engine document OCR pipeline (Tesseract, Surya OCR, Docling Layout, TableFormer) achieving a 40% improvement in table structure recognition, processing documents across 23 Indian languages.
- Developed a real-time TTS and translation web utility with optimized DOM traversal and request batching for seamless in-browser multilingual experiences.
- Built and deployed a Python REST backend integrating multiple AI/ML inference engines into a unified API layer serving 23 languages.

Software / Flutter Developer

SkoolSaver **November 2024 - July 2025**

- Architected a Node.js/Express REST API with dual-database architecture (PostgreSQL + MongoDB), OAuth-based JWT authentication, and 15+ RESTful endpoints.
- Built a mentor-booking system with slot-availability checking and conflict resolution, and a job portal with multi-parameter search.
- Containerized the full stack with Docker Compose, orchestrating 4 services for consistent dev/deploy environments.

Full Stack Web Developer - Committee Head

Directorate of Student Affairs, SRM IST, Chennai **September 2023 - October 2024**

- Designed and deployed systems serving 10,000+ students with data-processing workflows and analytics reporting.
- Built enterprise application integrations for large-scale event data and database management.
- Led technical initiatives and mentored 8+ developers on system design.

PROJECT

A.R.T.E.M.I.S

Self · github.com/Anshumanv28/A.R.T.E.M.I.S · January 2026 - Present

- Architected a production-grade multi-agent RAG system (Generative AI, Artificial Intelligence) using a LangGraph Supervisor pattern with 3 specialized agent workflows (search, ingestion, collection management) and lightweight LLM routing for cost reduction.
- Engineered 5 document chunking strategies supporting 4 file formats, backed by a Qdrant vector database for multi-collection retrieval, with citation-grounded answers and explicit missing-context flags to mitigate hallucination.
- Built an MCP server with least-privilege scopes per sub-agent for model security, and an automated evaluation harness tracking model accuracy, model reliability, model performance, and model cost across routing and tool-selection tasks; CI-tested via GitHub Actions, deployed via Docker + docker-compose with a live AWS demo.

AI_Concierge (Voice Booking Agent)

Self · github.com/Anshumanv28/AI_Concierge · December 2025 - Present

- Built a voice-native booking agent (LiveKit real-time voice + OpenAI GPT tool calling) and solved unreliable LLM function-calling for self-referential actions (call-ending, mic control) by isolating a second, narrow structured-decision LLM call instead of relying on prompting alone.
- Grounded booking logic in real computed alternative time slots and live weather data instead of LLM-invented responses, with centralized conflict-aware availability checks across every booking code path.
- Shipped a persistent chat-history dashboard, a custom mono-design frontend redesign (with accessibility fixes), and a 3-service CI/CD pipeline (GitHub Actions) pushing Docker images to GHCR on every merge.

EDUCATION

Bachelor of Technology

SRM Institute of Science and Technology • Chennai • 2021-25 • 8.67

SKILLS

Python, JavaScript, TypeScript, SQL

LLMs: OpenAI GPT, Prompt Engineering, Function/Tool Calling, Structured Outputs

Agentic AI: LangChain, LangGraph, Agent Workflows, Multi-Agent Orchestration, RAG Pipelines

Vector Search: Qdrant, Embeddings, Semantic Search

ML/NLP: PyTorch, ASR, TTS, Machine Translation, OCR (Tesseract, Docling, TableFormer)

Backend & Deployment: FastAPI, Flask, Node.js, REST APIs, Databases (PostgreSQL, MongoDB), Docker, AWS, Git/CI-CD